

1. Implement Polynomial regression on Data1.csv. Display the coefficients.

```
import pandas as pd
```

```
# Polynomial Regression Function
```

```
def polynomial_regression(x, y, degree):
```

```
    n = len(x)
```

```
    X = [[xi ** d for d in range(degree + 1)] for xi in x] #
```

```
Generate X matrix
```

```
    XT = list(zip(*X)) # Transpose of X
```

```
    XTX = [[sum(XT[i][j] * X[j][k] for j in range(n)) for k in  
range(degree + 1)] for i in range(degree + 1)] #  $X^T * X$ 
```

```
    XTY = [sum(XT[i][j] * y[j] for j in range(n)) for i in  
range(degree + 1)] #  $X^T * Y$ 
```

```
    # Solve for coefficients using Gaussian elimination
```

```
    coeff = [XTY[i] / XTX[i][i] for i in range(degree + 1)] #
```

```
Assuming diagonal dominance
```

```
    return coeff
```

```
# Load data from CSV file
```

```
data = pd.read_csv('/content/Data1.csv')
```

```
# Extract input (X1) and target (y) variables
```

```
x = data['X1'].values # Input variable
```

```
y = data['Y'].values # Target variable
```

```
# Degree of polynomial
```

```
degree = 2 # Adjust as needed
```

```
# Perform polynomial regression
```

```
coefficients = polynomial_regression(x, y, degree)
```

```
# Display results
print("Coefficients:", coefficients)
```

2. Apply Logistic Regression on Pima Indian Diabetes dataset to predict the output.

```
import numpy as np
import pandas as pd
```

```
# Sigmoid function
def sigmoid(z):
    return 1 / (1 + np.exp(-z))
```

```
# Logistic Regression
def logistic_regression(X, y, learning_rate, iterations):
    n, m = X.shape
    w = np.zeros(m) # Initialize weights
    b = 0           # Initialize bias

    for _ in range(iterations):
        # Compute linear combination and predictions
        linear_model = np.dot(X, w) + b
        predictions = sigmoid(linear_model)

        # Compute gradients
        dw = (1 / n) * np.dot(X.T, (predictions - y))
        db = (1 / n) * np.sum(predictions - y)

        # Update weights and bias
        w -= learning_rate * dw
        b -= learning_rate * db
```

```
return w, b
```

```
# Load Pima Indian Diabetes Dataset
```

```
data = pd.read_csv('/content/diabetes.csv')
```

```
# Feature selection and preprocessing
```

```
X = data.iloc[:, :-1].values # All columns except the last one
```

```
y = data.iloc[:, -1].values # Target column
```

```
# Normalize features
```

```
X = (X - X.mean(axis=0)) / X.std(axis=0)
```

```
# Train logistic regression
```

```
learning_rate = 0.01
```

```
iterations = 1000
```

```
weights, bias = logistic_regression(X, y, learning_rate,  
iterations)
```

```
print("Weights:", weights)
```

```
print("Bias:", bias)
```

3. Predict the Digits in Images Using a Logistic Regression Classifier in Python.

```
import numpy as np
```

```
# Sigmoid function
```

```
def sigmoid(z):
```

```
    return 1 / (1 + np.exp(-z))
```

```
# Logistic Regression Model
```

```
def logistic_regression(X, y, lr=0.01, epochs=1000):
```

```
    m, n = X.shape
```

```
weights = np.zeros(n) # Initialize weights
bias = 0
```

```
for _ in range(epochs):
    # Linear model
    z = np.dot(X, weights) + bias
    # Prediction
    y_pred = sigmoid(z)
    # Compute gradients
    dw = (1/m) * np.dot(X.T, (y_pred - y))
    db = (1/m) * np.sum(y_pred - y)
    # Update weights and bias
    weights -= lr * dw
    bias -= lr * db
```

```
return weights, bias
```

```
# Prediction function
def predict(X, weights, bias):
    z = np.dot(X, weights) + bias
    y_pred = sigmoid(z)
    return (y_pred > 0.5).astype(int)
```

```
# Example Dataset (MNIST digits simplified for binary
classification: digit 0 vs 1)
# Here we use small dummy data for simplicity
X_train = np.array([[1, 2], [2, 1], [2, 3], [3, 4], [4, 3]]) #
Features
y_train = np.array([0, 0, 1, 1, 1]) # Labels (binary
classification)
```

```
# Train the model
weights, bias = logistic_regression(X_train, y_train)
```

```
# Test the model
X_test = np.array([[1, 1], [4, 4]])
predictions = predict(X_test, weights, bias)
print("Predictions:", predictions)
```

4. Apply K-means clustering on the following data.

.x = [4, 5, 10, 4, 3, 11, 14, 6, 10, 12]

y = [21, 19, 24, 17, 16, 25, 24, 22, 21, 21]

```
def k_means(x, y, k, iterations):
    # Initialize centroids as the first `k` points
    centroids = [(x[i], y[i]) for i in range(k)]

    for _ in range(iterations):
        clusters = [[] for _ in range(k)]

        # Assign points to nearest centroid
        for i in range(len(x)):
            distances = [((x[i] - cx) ** 2 + (y[i] - cy) ** 2) for cx,
cy in centroids]
            cluster_index = distances.index(min(distances))
            clusters[cluster_index].append((x[i], y[i]))

        # Update centroids
        centroids = [(sum([p[0] for p in cluster]) / len(cluster),
sum([p[1] for p in cluster]) / len(cluster))
if cluster else centroids[i]
for i, cluster in enumerate(clusters)]

    return centroids, clusters

# Example Usage
```

```

x = [4, 5, 10, 4, 3, 11, 14, 6, 10, 12]
y = [21, 19, 24, 17, 16, 25, 24, 22, 21, 21]
k = 2
iterations = 10
centroids, clusters = k_means(x, y, k, iterations)
print("Centroids:", centroids)
print("Clusters:", clusters)

```

5. Perform Hierarchical Clustering on

Mall_Customers_data.csv. draw the dendrogram.

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

# Load Mall_Customers_data.csv
data = pd.read_csv('/content/Mall_Customers.csv')

# Extract relevant columns (e.g., Annual Income and Spending Score)
# Adjust column names based on your dataset
features = data[['Annual Income (k$)', 'Spending Score (1-100)']].values

# Compute Euclidean distance between two points
def euclidean_distance(p1, p2):
    return np.sqrt(np.sum((p1 - p2) ** 2))

# Perform hierarchical clustering
def hierarchical_clustering(data):
    n = len(data)
    clusters = {i: [i] for i in range(n)} # Each point is its own cluster

```

```

distances = [[euclidean_distance(data[i], data[j]) for j in
range(n)] for i in range(n)]
merges = []

while len(clusters) > 1:
    # Find closest clusters
    min_dist = float('inf')
    to_merge = None
    for i in clusters:
        for j in clusters:
            if i < j:
                dist = min(distances[p1][p2] for p1 in clusters[i]
for p2 in clusters[j])
                if dist < min_dist:
                    min_dist = dist
                    to_merge = (i, j)

    # Merge clusters
    i, j = to_merge
    merges.append((i, j, min_dist))
    clusters[i].extend(clusters[j])
    del clusters[j]

return merges

# Plot dendrogram
def plot_dendrogram(merges):
    plt.figure(figsize=(8, 5))
    current_positions = {i: i for i in range(len(merges) + 1)}
    for i, (a, b, height) in enumerate(merges):
        x1, x2 = current_positions[a], current_positions[b]
        x_mid = (x1 + x2) / 2
        plt.plot([x1, x1, x2, x2], [0, height, height, 0], 'b')

```

```
    current_positions[a] = x_mid # Update cluster position
    del current_positions[b]
plt.xlabel("Data Points")
plt.ylabel("Distance")
plt.title("Dendrogram")
plt.show()
```

```
# Run and plot
merges = hierarchical_clustering(features)
plot_dendrogram(merges)
```

7.Find the optimal hyperplane for SVM use the following data set

positive class:(3,1),(3,-1),(6,1),(6,-1)

Negative Class:(1,0),(0,1),(0,-1),(-1,0)

```
import numpy as np
```

```
# Dataset: Positive and Negative classes
```

```
positive_class = np.array([[3, 1], [3, -1], [6, 1], [6, -1]])
```

```
negative_class = np.array([[1, 0], [0, 1], [0, -1], [-1, 0]])
```

```
# Combine data and labels
```

```
X = np.vstack((positive_class, negative_class))
```

```
y = np.hstack((np.ones(len(positive_class)),  
-np.ones(len(negative_class))))
```

```
# Helper functions for SVM
```

```
def compute_svm(X, y):
```

```
    n_samples, n_features = X.shape
```



```
# Initialize weights and bias
```

```
w = np.zeros(n_features)
```

```
b = 0
```

```
lr = 0.01 # Learning rate
```

```
epochs = 1000
```

```
# Gradient Descent for optimization
```

```
for _ in range(epochs):
```

```
    for i in range(n_samples):
```

```
        if y[i] * (np.dot(w, X[i]) + b) < 1: # Misclassified
```

```
            w += lr * (y[i] * X[i] - 2 * (1 / epochs) * w)
```

```
            b += lr * y[i]
```

```
        else: # Correct classification
```

```
            w -= lr * (2 * (1 / epochs) * w)
```

```
    return w, b
```

```
# Solve for the optimal hyperplane
```

```
w, b = compute_svm(X, y)
```

```
# Decision boundary equation
```

```
print(f"Optimal weight vector: {w}")
```

```
print(f"Optimal bias: {b}")
```

```
print(f"Decision boundary: {w[0]} * x1 + {w[1]} * x2 + {b} = 0")
```

```
# Plot the data and decision boundary
```

```
import matplotlib.pyplot as plt
```

```
# Plot data points
```

```
plt.scatter(positive_class[:, 0], positive_class[:, 1],
```

```
color='blue', label='Positive Class')
```

```
plt.scatter(negative_class[:, 0], negative_class[:, 1],  
color='red', label='Negative Class')
```

```
# Plot decision boundary
```

```
x1 = np.linspace(-2, 7, 100)
```

```
x2 = -(w[0] * x1 + b) / w[1]
```

```
plt.plot(x1, x2, color='green', label='Decision Boundary')
```

```
plt.xlabel('x1')
```

```
plt.ylabel('x2')
```

```
plt.title('SVM Decision Boundary')
```

```
plt.legend()
```

```
plt.grid()
```

```
plt.show()
```